

Inference of Natural Selection from Interspersed Genomic Elements Based on Polymorphism and Divergence

Ilan Gronau,¹ Leonardo Arbiza,¹ Jaaved Mohammed,^{1,2} and Adam Siepel^{*1}

¹Department of Biological Statistics and Computational Biology, Cornell University

²Tri-Institutional Training Program in Computational Biology and Medicine, New York, NY

***Corresponding author:** E-mail: acs4@cornell.edu.

Associate editor: Daniel Falush

Abstract

Complete genome sequences contain valuable information about natural selection, but this information is difficult to access for short, widely scattered noncoding elements such as transcription factor binding sites or small noncoding RNAs. Here, we introduce a new computational method, called *Inference of Natural Selection from Interspersed Genomically coHerent elemenTs* (INSIGHT), for measuring the influence of natural selection on such elements. INSIGHT uses a generative probabilistic model to contrast patterns of polymorphism and divergence in the elements of interest with those in flanking neutral sites, pooling weak information from many short elements in a manner that accounts for variation among loci in mutation rates and coalescent times. The method is able to disentangle the contributions of weak negative, strong negative, and positive selection based on their distinct effects on patterns of polymorphism and divergence. It obtains information about divergence from multiple outgroup genomes using a general statistical phylogenetic approach. The INSIGHT model is efficiently fitted to genome-wide data using an approximate expectation maximization algorithm. Using simulations, we show that the method can accurately estimate the parameters of interest even in complex demographic scenarios, and that it significantly improves on methods based on summary statistics describing polymorphism and divergence. To demonstrate the usefulness of INSIGHT, we apply it to several classes of human noncoding RNAs and to GATA2-binding sites in the human genome.

Key words: molecular evolution, population genetics, noncoding DNA, regulatory sequences, probabilistic graphical models.

Introduction

Evolutionary modeling has become an essential tool in genomic analysis, particularly for the study of noncoding elements in eukaryotes, which tend to be sparsely annotated, poorly understood, and difficult to examine experimentally. So far, most evolutionary analyses of such elements have been based on patterns of sequence divergence between genomes that diverged millions of years ago (Boffelli et al. 2003; Thomas et al. 2003; Cooper et al. 2005; Siepel et al. 2005; Pollard et al. 2010). However, numerous confounding factors limit the utility of this approach. For example, evolutionary “turn-over” (gain and loss of functional elements) can be prominent on these time scales (Dermitzakis and Clark 2002; Moses et al. 2006) and can distort patterns of sequence divergence. In addition, positive and negative selection can sometimes act on the same sequences and have partially canceling effects on divergence. Finally, technical challenges such as orthology identification and genomic alignment are nontrivial on these time scales, and errors in these procedures can produce spurious inferences of natural selection.

In principle, data describing genetic polymorphism within species could help to address these limitations. Patterns of polymorphism reflect evolutionary processes on relatively short timescales, during which turnover should be much less prevalent. Orthology identification and alignment are

also much more straightforward on these time scales. Furthermore, it is well known that patterns of polymorphism within a species and divergence between species can be used to tease apart the signatures of positive and negative selection (McDonald and Kreitman 1991; Sawyer and Hartl 1992; Bustamante et al. 2005).

In practice, however, it is technically challenging to extract useful information about natural selection from patterns of polymorphism and divergence in noncoding elements. Many of the elements of interest, such as transcription factor binding sites and small noncoding RNAs, are quite short (at most tens of bases in length), and polymorphisms tend to be sparse. As a result, many elements of interest typically contain no informative sites whatsoever, whereas most others contain just one or two. This problem can be addressed by pooling data from multiple elements (Andolfatto 2005), but variation across loci in mutation rates and coalescence times can lead to difficulties in interpreting such pooled data sets (Smith and Eyre-Walker 2002; Stoletzki and Eyre-Walker 2011). Finally, the confounding influence of demography on patterns of polymorphism is a persistent problem when attempting to draw conclusions about natural selection (Nielsen et al. 2007).

Here, we describe a new computational method, called *Inference of Natural Selection from Interspersed Genomically coHerent elemenTs* (INSIGHT), that is designed

to address these challenges. INSIGHT uses the general strategy of contrasting patterns of polymorphism and divergence in a collection of elements of interest with those in flanking neutral regions, thereby mitigating biases from demography, variation in mutation rates, and differences in coalescence time. In this way, it resembles McDonald–Kreitman (MK)-based methods for identifying departures from neutrality (McDonald and Kreitman 1991; Sawyer and Hartl 1992; Smith and Eyre-Walker 2002; Andolfatto 2005). Unlike these methods, however, INSIGHT is based on a generative probabilistic model, accommodates weak negative (WN) selection (Charlesworth and Eyre-Walker 2008), and allows diffuse information from many short elements across the genome to be pooled efficiently, in a manner that avoids statistical pitfalls arising from pooling counts of site classes. Our modeling approach also fully integrates phylogenetic information from multiple outgroup species with genome-wide population genetic data.

The main purpose of this article is to detail the probabilistic model and inference strategy underlying the INSIGHT method. We also compare our full probabilistic model with summary-statistic-based methods similar to those used in a number of previous polymorphism-and-divergence studies (Fay et al. 2001; Smith and Eyre-Walker 2002; Andolfatto 2005), demonstrating several advantages of our methods across a range of simulation parameters. In a parallel submission (Arbiza L, Gronau I, Aksoy BA, Hubisz MJ, Gulko B, Keinan A, Siepel A, in revision), we report the use of INSIGHT in a large-scale analysis of transcription factor binding sites in the human genome, based on

chromatin-immunoprecipitation-and-sequencing (ChIP-seq) data for 78 human transcription factors (Dunham et al. 2012). Here, we further demonstrate the breadth of applicability of the method by applying it to several classes of human noncoding RNAs and using it to carry out a position-specific analysis of GATA2-binding sites in the human genome.

Materials and Methods

General Approach

The central goal of INSIGHT is to characterize the aggregate influence of natural selection on a collection of elements having some arbitrary genomic distribution (fig. 1). The collection of elements is assumed to be reasonably homogeneous and coherent but can be defined in many different ways. For example, it might include all binding sites of a particular transcription factor, all noncoding RNAs of a particular type, all binding sites near genes of a particular functional category, or all paired bases in a group of RNAs (see Discussion). We assume the individual elements are fairly short, ranging from a single nucleotide to perhaps a few hundred bases in length. The key modeling challenge is to integrate sparse information from many such elements in a manner that accounts for variation along the genome in properties such as mutation rate and coalescence time. Rather than attempting to fully describe the relationships among selection, polymorphism, and divergence—which is complex and demography-dependent—our model works by contrasting patterns of polymorphism and divergence in the elements of interest with those in nearby neutral sites.

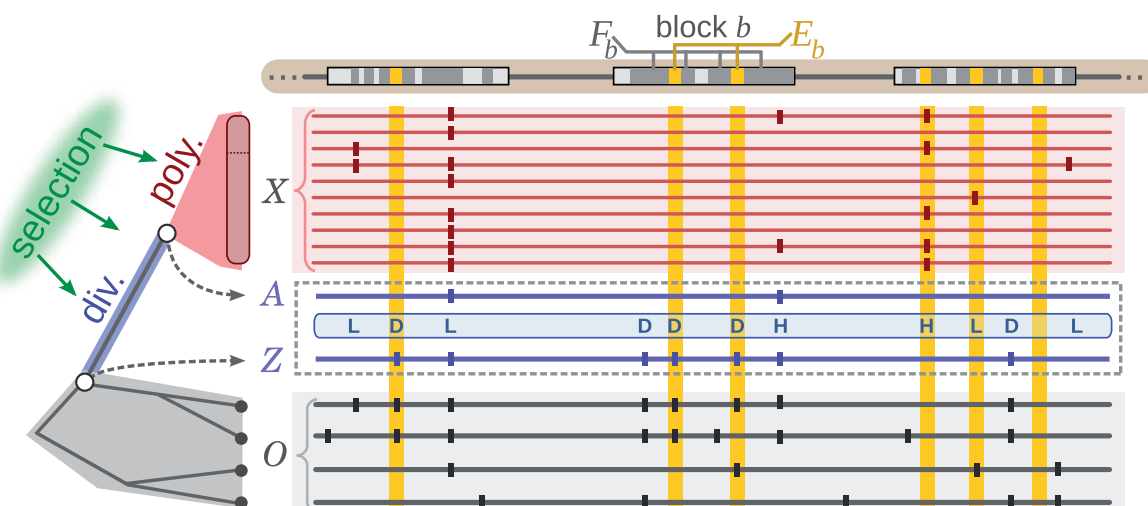


FIG. 1. Schematic description of INSIGHT. The method measures the influence of natural selection by contrasting patterns of polymorphism and divergence in a collection of genomic elements of interest (gold) with those in flanking neutral sites (dark gray). Nucleotide sites in both elements (E_b) and flanks (F_b) are grouped into genomic blocks of a few kilobases in length (b) to accommodate variation along the genome in mutation rate and coalescence time. The model consists of phylogenetic (gray), recent divergence (blue), and intraspecies polymorphism (red) components, which are applied to genome sequences for the target population (X , red) and outgroup species (O , gray). At each nucleotide position, the alleles at the MRCA of the samples from the target population (A) and of the target population and closest outgroup (Z) are represented as hidden variables and treated probabilistically during inference. The allele Z determines whether monomorphic sites are considered to be divergent (D). Polymorphic sites are classified as having low- (L) or high- (H) frequency-derived alleles based on A and a frequency threshold f . The labels shown here are based on a likely setting of Z and A . Vertical ticks represent single nucleotide variants relative to an arbitrary reference. Inference is based on differences in the patterns of polymorphism and divergence expected at neutral and selected sites.

We assume that genome-wide polymorphism data are available for a particular target population, in a form that allows polymorphic sites to be reliably distinguished from invariant sites and provides reasonably accurate information about allele frequencies. At present, this is most easily achieved using high-coverage individual genome sequences, although our methods could also be adapted to make use of statistically inferred genotype frequencies based on low-coverage sequence data (Yi et al. 2010). We further assume that genomic sequence data are available for one or more outgroup species. Although the method can be used with a single outgroup genome, better information about ancestral alleles can be obtained by using two or more minimally distant outgroups that diverged from one another prior to the divergence of either one from the target population.

We assume that each nucleotide site evolves according to one of four possible selective modes: neutral drift (neut), strong negative (SN) selection, WN selection, or positive selection (P). This coarse-grained, categorical approach to modeling the distribution of fitness effects (DFE) is motivated by observations indicating that the data contain only limited information about the full DFE (Boyko et al. 2008; Wilson et al. 2011). The key to our approach is that these four selective modes have qualitatively distinct effects on patterns of polymorphism and divergence (cf., Bierne and Eyre-Walker 2004). In particular, SN and positive selection will generally cause mutations to reach fixation or be lost rapidly, and therefore will mostly eliminate observable polymorphisms. By contrast, WN selection will allow polymorphisms to persist for longer periods of time, but will tend to hold derived alleles at low frequencies. In addition, negative selection (either strong or weak) will largely prohibit the eventual fixation of derived alleles. Therefore, we make the following three assumptions about nucleotide sites under selection: 1) only positively selected sites make nonnegligible contributions to divergence; 2) only WN sites make nonnegligible contributions to polymorphism; and 3) any polymorphisms must have low derived allele frequencies. (Neutral sites, of course, may also contribute to divergence and polymorphism.) Together, these assumptions allow the fraction of sites under selection to be estimated. As it turns out, they are not sufficient to fully disentangle the contributions of all four selective modes, but they do allow us to obtain indirect information about

the contributions of positive and WN selection at selected sites (discussed later).

In addition, we classify every site as monomorphic (M), polymorphic with a low-frequency-derived allele (L), or polymorphic with a high-frequency-derived allele (H), where the distinction between L and H sites depends on a designated low-frequency threshold f (typically $f = 0.15$). Information about selection comes from the relative frequencies of these labels in the elements of interest relative to the flanking neutral sites, together with patterns of divergence with respect to the outgroup genomes. A minor complication is that in some cases, the derived allele class depends on the ancestral allele, which is not known. We address this problem by treating the ancestral allele as a hidden (latent) random variable and integrating over possible values as needed. The use of this low-dimensional projection of the SFS is intended to buffer our method from the effects of recent demographic changes in the target population. In the simulation analyses reported later, we examine the extent to which our inferences are robust to demography. We also examine their dependence on the threshold f .

Probabilistic Model

Our model assumes that the genomic regions under study are partitioned into a collection of blocks, B . The nucleotide sites within each block $b \in B$ are further partitioned into sites within the elements of interest, E_b , and the associated neutral flanking sites, F_b (cumulatively E and F , respectively). Each block is assigned a population-scaled mutation rate (θ_b), a neutral divergence scale factor (λ_b), and an outgroup divergence scale factor (λ_b^O). In addition, the model has four global parameters: the fraction of sites under selection in elements (ρ), the relative divergence (η) and polymorphism (γ) rates at selected sites, and β , a multivariate parameter summarizing the neutral site frequency spectrum (table 1). The full set of parameters is denoted by ζ .

Each site i is associated with a set of aligned bases from the outgroup genomes (O_i) and the polymorphism data for the target population (X_i). X_i is further summarized as $X_i = (X_i^{\text{maj}}, X_i^{\text{min}}, Y_i)$, where X_i^{maj} and X_i^{min} are the observed major and minor alleles, and $Y_i \in \{M, L, H\}$ is the minor allele frequency class ($X_i^{\text{min}} = \emptyset$ when $Y_i = M$). The entire data set is denoted by (\mathbf{X}, \mathbf{O}) . Y_i is defined by the observed minor allele frequency m_i and the specified low-frequency threshold,

Table 1. Model Parameters.

Parameter	Type	Description
$\lambda^O = \{\lambda_b^O\}_{b \in B}$	Neutral	Block-specific neutral scaling factor for the outgroup portion of the phylogeny, used when computing the prior distributions for the deep ancestral allele $P(Z_i O_i, \lambda_b^O)$.
$\lambda = \{\lambda_b\}_{b \in B}$	Neutral	Block-specific neutral scaling factor for divergence.
$\theta = \{\theta_b\}_{b \in B}$	Neutral	Block-specific neutral polymorphism rate.
$\beta = (\beta_1, \beta_2, \beta_3)$	Neutral	Relative frequencies of the three derived allele frequency classes, $(0, f)$, $[f, 1 - f]$, and $(1 - f, 1)$, within neutral polymorphic sites.
ρ	Selection	Fraction of sites under selection within functional elements.
η	Selection	Ratio of divergence rate at selected sites to local neutral divergence rate.
γ	Selection	Ratio of polymorphism rate at selected sites to local neutral polymorphism rate.

Table 2. Model Variables Associated with Site i .

Variable	Type	Description
O_i	Observed	Set of aligned bases from outgroup species
X_i^{maj}	Observed	Base for major allele in target population
X_i^{min}	Observed	Base for minor allele in target population (\emptyset for monomorphic sites)
Y_i	Observed	MAF class for site i : “M” for monomorphic sites (MAF = 0) “L” for polymorphic sites with MAF < f “H” for polymorphic sites with MAF $\geq f$
S_i	Hidden	Selection class: “neut” for neutral sites “sel” for sites under selection
Z_i	Hidden	Ancestral allele at the MRCA of the target population and the closest outgroup
A_i	Hidden	Ancestral allele at the MRCA of samples from the target population

$f < \frac{1}{2}$; in particular, $Y_i = M$ when $m_i = 0$, $Y_i = L$ when $0 < m_i < f$, and $Y_i = H$ when $m_i \geq f$. Note that the minor allele frequency is observed, whereas the derived allele frequency depends on the identity of the hidden population ancestral allele (A_i) and is thus inferred probabilistically by the model. Sites with three or more alleles are discarded in preprocessing. Each site is also associated with three hidden variables: a selection class ($S_i \in \{\text{sel}, \text{neut}\}$), a “deep” ancestral allele at the most recent common ancestor (MRCA) of the target population and closest outgroup (Z_i), and a population ancestral allele (A_i) (table 2).

We assume independence of blocks, conditional independence of the nucleotide sites within each block given the model parameters, conditional independence of the variables describing the target population (A_i , S_i , and X_i) from the outgroups (O_i) given the deep ancestral allele (Z_i), and independence of the site-wise selection classes (S_i), as shown graphically in figure 2. The same graphical model applies to all sites, except that the selection class is fixed to “neut” for the flanking sites. Thus, a likelihood function for the model, conditional on the outgroup data, can be written as follows:

$$\begin{aligned} \mathcal{L}(\zeta; \mathbf{X}, \mathbf{O}) &\equiv P(\mathbf{X} | \mathbf{O}, \zeta) \\ &= \prod_{b \in B} \left[\prod_{i \in F_b} \sum_z \sum_a P(X_i, Z_i = z, A_i = a | S_i = \text{neut}, O_i, \zeta) \right] \\ &\quad \times \left[\prod_{i \in F_b} \sum_{s \in \{\text{neut}, \text{sel}\}} P(S_i = s | \zeta) \sum_z \sum_a P(X_i, Z_i = z, A_i = a | S_i = s, O_i, \zeta) \right]. \end{aligned} \quad (1)$$

Furthermore, each term of the form $P(X_i, Z_i, A_i | S_i, O_i, \zeta)$ can be factorized as follows:

$$\begin{aligned} P(X_i, Z_i, A_i | S_i, O_i, \zeta) \\ = P(Z_i | O_i, \lambda_b^O) P(A_i | S_i, Z_i, \zeta) P(X_i | S_i, A_i, Z_i, \zeta). \end{aligned} \quad (2)$$

This likelihood function is composed of four conditional probability distributions, corresponding to the variables S_i , Z_i , A_i , and X_i . The distribution for S_i is needed only for element sites and is given by a two-component mixture model with coefficient ρ :

$$P(S_i = s | \zeta) = \begin{cases} \rho & s = \text{sel} \\ 1 - \rho & s = \text{neut} \end{cases}. \quad (3)$$




Fig. 2. Graphical model for an individual nucleotide site i . As in figure 1, the phylogenetic portion of the model is shown in gray, the divergence component in blue, and the polymorphism component in red. Observed variables are represented by solid circles and hidden variables by empty circles. The observed alleles in the target population and outgroups are represented by X_i and O_i , respectively. X_i consists of a major (X_i^{maj}) and minor (X_i^{min}) allele, as well as the minor allele frequency class (Y_i ; not shown). The selection class is denoted S_i , and the ancestral alleles are denoted Z_i and A_i , as described in figure 1. Conditional dependence between the variables is indicated by directed edges, in the standard manner for probabilistic graphical models. Model parameters are shown alongside the associated conditional dependency edges. The selection parameters $\zeta_{\text{sel}} = (\rho, \eta, \gamma)$ are highlighted in green.

The conditional distribution for Z_i given the outgroup data, $P(Z_i | O_i, \lambda_b^O)$, is based on a standard statistical phylogenetic model and is computed using existing software. Notice that our model assumes that the phylogenetic model for the outgroups is independent of the selection class, S_i . This assumption is not strictly warranted (sites under selection are likely to evolve at different rates in the outgroups), but it dramatically simplifies the inference procedure by allowing us to pre-estimate the outgroup scale factors (λ_b^O) and the site-wise distributions for Z_i (see Parameter Inference). In practice, this simplifying assumption is of little consequence, because it only affects the prior distribution for Z_i , which is fairly insensitive to evolutionary rates in outgroup lineages as long as the branches of the phylogeny are not too long.

The third conditional distribution, $P(A_i | S_i, Z_i, \zeta)$, describes the process of sequence divergence on the lineage leading to the target population. Given a global neutral branch length t for this lineage (in substitutions per site),

we assume a nucleotide substitution rate of $\lambda_b t$ for neutral sites and $\eta \lambda_b t$ for sites under selection. Note that η can be driven downward by negative selection or upward by positive selection, so that it may be greater or less than one, depending on the DFE. Because we are primarily interested in cases in which t is quite small (e.g., $t \approx 0.005$ for the case of humans and chimpanzees), we use the following approximation for the probability of divergences under a Poisson substitution model:

$$P(A_i = a \mid S_i = s, Z_i = z, \zeta) = \begin{cases} \frac{1}{3} \lambda_b t & s = \text{neut}, a \neq z \\ 1 - \lambda_b t & s = \text{neut}, a = z \\ \frac{1}{3} \eta \lambda_b t & s = \text{sel}, a \neq z \\ 1 - \eta \lambda_b t & s = \text{sel}, a = z. \end{cases} \quad (4)$$

Finally, the fourth conditional distribution, $P(X_i \mid S_i, A_i, Z_i, \zeta)$, describes the patterns of polymorphism in the target population given the ancestral alleles and selection class. The definition of this distribution is somewhat more involved. Briefly, we assume that the neutral polymorphism rate for each genomic block b is determined by a block-specific population-scaled mutation rate, $\theta_b = 4N_b \mu_b$, which allows the model to accommodate both variable mutation rates and selection from linked sites (background selection or hitchhiking). The probability of observing a polymorphic nucleotide position is taken to be $\theta_b a_n$, where n is the number of haploid genomes sampled and $a_n = \sum_{k=1}^{n-1} 1/k$ (Watterson 1975). In the absence of missing data, a_n is a constant of no consequence in the inference procedure, but it can be used to accommodate sites with small amounts of missing genotype data if desired (see Discussion and [supplementary methods, Supplementary Material](#) online). Given a neutral polymorphism, the probabilities of low-, intermediate-, and high-frequency-derived allele are given by β_1 , β_2 , and β_3 , respectively ($\sum_{i=1}^3 \beta_i = 1$). The situation is similar for sites under selection, except that they are assumed to have population-scaled mutation rates of $\gamma \theta_b$ and only low-frequency derived alleles are permitted. It is possible to derive closed-form expressions for this distribution for all cases of interest ([supplementary table S1, Supplementary Material](#) online). In addition, the conditional distributions $P(X_i \mid S_i, A_i, Z_i, \zeta)$ and $P(A_i \mid S_i, Z_i, \zeta)$ can be combined into a single conditional distribution table, $P(X_i \mid S_i, Z_i, \zeta)$, by integrating over possible values of A_i . This integration is simplified by assuming an infinite sites model for the time since the population-level MRCA, which implies $A_i \in \{X_i^{\text{maj}}, X_i^{\text{min}}\}$ ([table 3](#)).

Parameter Inference

The main objective of the inference procedure is to produce maximum likelihood estimates (MLEs) of the selection parameters, ρ , η , and γ , but to do so, the neutral parameters $\zeta_{\text{neut}} = (\lambda^0, \lambda, \theta, \beta)$ must also be estimated. In principle, an expectation-maximization (EM) algorithm could be used to jointly estimate all model parameters. However, this approach is impractical for genome-wide applications involving millions

Table 3. Conditional Distribution Table for $P(X_i \mid S_i, Z_i, \zeta)$.

S	y	z, $x^{\text{maj}}, x^{\text{min}}$ ^a	$P(X_i = (x^{\text{maj}}, x^{\text{min}}, y) \mid S_i = s, Z_i = z, \zeta)$
neut	M	$z = x^{\text{maj}}$	$(1 - \lambda_b t)(1 - \theta_b a_n)$
neut	M	$z \neq x^{\text{maj}}$	$\frac{1}{3} \lambda_b t (1 - \theta_b a_n)$
neut	L	$z = x^{\text{maj}}$	$((1 - \lambda_b t)\beta_1 + \frac{1}{3} \lambda_b t \beta_3) \frac{1}{3} \theta_b a_n$
neut	L	$z = x^{\text{min}}$	$((1 - \lambda_b t)\beta_3 + \frac{1}{3} \lambda_b t \beta_1) \frac{1}{3} \theta_b a_n$
neut	L	$z \notin \{x^{\text{maj}}, x^{\text{min}}\}$	$\frac{1}{3} \lambda_b t (\beta_1 + \beta_3) \frac{1}{3} \theta_b a_n$
neut	H	$z \in \{x^{\text{maj}}, x^{\text{min}}\}$	$(1 - \lambda_b t + \frac{1}{3} \lambda_b t) \beta_2 \frac{1}{3} \theta_b a_n$
neut	H	$z \notin \{x^{\text{maj}}, x^{\text{min}}\}$	$\frac{2}{3} \lambda_b t \beta_2 \frac{1}{3} \theta_b a_n$
sel	M	$z = x^{\text{maj}}$	$(1 - \eta \lambda_b t)(1 - \gamma \theta_b a_n)$
sel	M	$z \neq x^{\text{maj}}$	$\frac{1}{3} \eta \lambda_b t$
sel	L	$z = x^{\text{maj}}$	$(1 - \eta \lambda_b t) \frac{1}{3} \gamma \theta_b a_n$
sel	L	$z \neq x^{\text{maj}}$	0
sel	H	—	0

^aRelationships among variables. It is implicit that $x^{\text{maj}} \in \{A, C, G, T\}$ and $x^{\text{maj}} \neq x^{\text{min}}$ in all cases. In addition, $x^{\text{min}} = \emptyset$ when $y = M$.

of nucleotide sites. Instead, we take advantage of the “loose coupling” between the phylogenetic outgroup model and the remaining portions of the model, and between the portions of the model concerned with the elements and the flanking sites, to decompose the inference procedure into separate stages, each of which can be performed fairly simply and efficiently.

Our inference procedure is based on the observation that the likelihood function can be expressed as a product of a function of the flanking sites and a function of the element sites (eq. 1). The first function depends only on the neutral parameters, whereas the second function depends on both the neutral and the selection parameters. However, because the flanking sites are expected to significantly outnumber the sites within the elements, the information about the neutral parameters comes predominately from the first function, and they can be estimated to a good approximation by maximizing this function only. The selection parameters can then be estimated by conditionally maximizing the second function. By making some additional minor simplifying assumptions, the first stage of inference can be further divided into two separate steps, one concerned with estimation of the phylogenetic parameters, λ^0 and λ , and one concerned with estimation of the population genetic parameters, θ and β . Our inference procedure thus consists of the following three distinct stages ([supplementary methods, Supplementary Material](#) online):

Phylogenetic Model Fitting: The divergence scale factors λ_b and λ_b^0 are estimated by fitting a pre-estimated neutral phylogenetic model to putative neutral sites in each genomic block using standard phylogenetic fitting procedures (Hubisz et al. 2011). The fitted phylogenetic model for the outgroup species is then used to compute the prior distribution for ancestral alleles, $P(Z_i \mid O_i, \lambda_b^0)$, at all sites in the block.

Neutral Polymorphism Model Fitting: MLEs of the block-specific polymorphism rate parameters, θ_b , and the global parameter β_2 are obtained using simple closed-form expressions. Global parameters β_1 and β_3

are estimated by a simple EM algorithm (they do not have closed-form estimators due to ancestral uncertainty).

Selection Inference: The selection parameters ρ , η , and γ are estimated conditional on the pre-estimated neutral parameters and ancestral priors by maximizing the likelihood of the element sites only by EM.

Extracting Information about the Modes of Selection

Although the INSIGHT model does not permit direct estimation of the fractions of sites under WN, SN, or positive selection, the estimated model parameters can be used to obtain indirect measures of the impact of WN and positive selection. A useful measure of positive selection is D_p , the number of divergence events driven by positive selection (sometimes called “adaptive substitutions”), and a similar measure pertaining to WN selection is P_w , the number of polymorphic sites subject to selection. Expected values for D_p and P_w can be obtained by summing over site-wise posterior probabilities associated with the variable configurations ($Y_i = M, Z_i \neq A_i, S_i = \text{sel}$) and ($Y_i = L, S_i = \text{sel}$), respectively (supplementary methods, Supplementary Material online). To allow comparisons between sets of different sizes, we normalize $\mathbb{E}[D_p]$ and $\mathbb{E}[P_w]$ by dividing them by the total number of nucleotide sites considered (in kilobases). By dividing $\mathbb{E}[D_p]$ by the total (expected) number of divergences, one can alternatively obtain an estimate of the fraction of substitutions driven by positive selection, a quantity known as α (Smith and Eyre-Walker 2002; Andolfatto 2005) (supplementary methods, Supplementary Material online).

Confidence Intervals and Likelihood Ratio Tests

The probabilistic nature of the model allows us to estimate standard errors for the estimated selection parameters using the curvature method (Lehmann and Casella 1998), based on an approximate Fisher information matrix derived from the 3×3 matrix of second derivatives for the log-likelihood function for ρ , η , and γ (eq. 1) at the joint MLE (supplementary methods, Supplementary Material online). In addition, likelihood ratio tests (LRTs) can be used to evaluate evidence for selection in general ($\rho > 0$), positive selection ($\eta > 0$), and WN selection ($\gamma > 0$). The LRTs are performed by fitting the model to the data twice, once with no restrictions on the free parameters, and once with a parameter of interest fixed at zero. Twice the difference in log likelihoods is then treated as a test statistic and compared with an appropriate asymptotic distribution. The tests for $\eta > 0$ and $\gamma > 0$ involve nested models in which the null hypothesis falls at a boundary of the alternative hypothesis. The associated test statistics therefore have asymptotic null distributions equal to a 50:50 mixture of a χ^2 distribution with one degree of freedom and a point mass at zero (Chernoff 1954; Self and Liang 1987). The case of ρ is more complex, because a value of $\rho = 0$ causes η and γ to become irrelevant to the likelihood function. Therefore, we used an empirical distribution to determine the cutoff for this LRT and found this to be consistent with a χ^2 distribution with three degrees of freedom (see Results).

Implementation and Software

The INSIGHT software consists of several modules. The main module is a C program, INSIGHT-EM, for the two EM algorithms used for inference: the main one for the selection parameters and a simpler one for β_1 and β_3 . The phylogenetic model fitting stage is implemented separately using procedures from RPHAST (Hubisz et al. 2011), and additional scripts are used for processing and filtering the polymorphism data. The INSIGHT website (<http://compngen.bscb.cornell.edu/INSIGHT/>, last accessed February 15, 2013) provides source code, documentation, and sample files for running the EM algorithm. The website also provides access to a server that can be used to run INSIGHT on any collection of human genomic elements, using our precomputed summaries of human polymorphism data (discussed later).

Estimators Based on Summary Statistics

For comparison with our model-based estimates, we made use of simple estimators for the fraction of sites under selection (ρ) and the number of adaptive substitutions (D_p). These estimators are based on the numbers of polymorphisms in element and flanking sites, denoted P_E and P_F , respectively, and the numbers of divergence events in element and flanking sites, denoted D_E and D_F , respectively. They include a divergence-based estimator for ρ introduced by Kondrashov and Crow (1993),

$$\hat{\rho}_{\text{Div}} = 1 - \frac{D_E |F|}{|E| D_F}, \quad (5)$$

a parallel estimator based on polymorphism rates,

$$\hat{\rho}_{\text{Poly}} = 1 - \frac{P_E |F|}{|E| P_F}, \quad (6)$$

and an estimator for $\mathbb{E}[D_p]$ based on the McDonald and Kreitman (1991) test, adapted from Smith and Eyre-Walker (2002):

$$\hat{D}_{p\text{-MK}} = D_E - \frac{P_E D_F}{P_F}. \quad (7)$$

In comparison with our model-based estimates, the divergence-based estimator $\hat{\rho}_{\text{Div}}$ ignores the effect of positive selection, and the estimators $\hat{\rho}_{\text{Poly}}$ and $\hat{D}_{p\text{-MK}}$ both implicitly assume that no polymorphisms occur in selected sites, and thus ignore the effects of WN selection. All three estimators share the limitation of pooling counts across elements in a manner that does not account for variable mutation rates across loci.

Simulations

We conducted a series of experiments on simulated data to assess the validity of our modeling assumptions and to evaluate the accuracy of the inference method. Simulated elements and flanking regions were generated with the forward simulator SFS_CODE (Hernandez 2008), assuming various mixtures of selective modes for the elements. We simulated data for human populations and chimpanzee, orangutan, and

rhesus macaque outgroups, using parameters based on previous studies. Each simulated block consisted of a 10 bp element, reflecting a typical binding site, and 5,000 flanking neutral sites on each side. We assumed a constant recombination rate and a randomly varying mutation rate, and each nucleotide position was assigned to one of four selection classes: neutral evolution ($2N_e s = 0$), SN selection ($2N_e s = -100$), WN selection ($2N_e s = -10$), and positive selection ($2N_e s = 10$). The choices of population-scaled selection coefficients were approximately based on several other recent studies (Eyre-Walker et al. 2006; Boyko et al. 2008; Wilson et al. 2011). Selection at WN and SN sites was held constant across the phylogeny, whereas for P sites, we assumed an interval of positive selection followed by WN selection on the lineage leading to the human population, to simulate selective sweeps rather than recurrent positive selection (see [supplementary methods, Supplementary Material](#) online, for complete details). The 10 kb flanking sites were all assigned to the neutral class, and the 10 bp of each simulated element were allocated among the four classes by multinomial sampling. When the simulation was done, sequence data were extracted using a single haploid sample from each of the three outgroup populations and 50 diploid samples from the target (human) population. In addition to assuming a range of mixtures of selective modes, we considered collections with various numbers of elements (ranging from 10,000 to 20,000), examined four different demographic scenarios, and perturbed the selection coefficients used for each category of selection (see [supplementary methods, Supplementary Material](#) online, for complete details).

The values of ρ , $\mathbb{E}[D_p]$ and $\mathbb{E}[P_w]$ estimated by INSIGHT were compared with “true” values for each simulation. The true value of ρ was simply the fraction of sites assumed to be under selection during data generation. The true value of D_p was taken to be the number of actual divergence events that occurred in sites under positive selection. The true value of P_w was taken to be the number of negatively selected sites that are polymorphic. In computing this quantity, we allowed for both strong and WN selection, because we are interested in accounting for all segregating deleterious alleles, regardless of our modeling assumptions. For ρ and D_p , we also compared our model-based estimates with the simple estimates based on counts of polymorphic and divergent sites (eq. 5–7).

Analysis of Human Noncoding Genomic Elements

In our analysis of real data, we made use of the 69 individual human genome sequences recently released by Complete Genomics (<http://www.completegenomics.com/public-data/69-Genomes/>, last accessed February 15, 2013) (Drmanac et al. 2010), using data for 54 unrelated individuals. Although larger data sets are available (1000 Genomes Project Consortium 2010), this one was selected for its high coverage, which reduces the effect of genotyping error and allows singleton variants to be characterized with fairly high confidence. For outgroup genomes, we used the chimpanzee (panTro2), orangutan (ponAbe2), and rhesus Macaque (rheMac2) reference genomes. Various filters were applied

to guarantee high quality alignments and variant calls ([supplementary methods, Supplementary Material](#) online). Putatively neutral sites were identified by excluding exons of known protein-coding and RNA genes plus 1 kb of flanking sites on each side, and previously predicted conserved noncoding elements plus flanking regions of 100 bp. After these filters were applied, an average of 3,881 sites per 10,000 bp block remained. Genomic blocks with <100 putative neutral sites were discarded.

We examined several classes of short interspersed noncoding elements in the human genome, including several collections of regulatory noncoding RNAs and a collection of GATA2 transcription factor binding sites. Annotations for noncoding RNAs were taken from GENCODE v.13 (Harrow et al. 2012) ([supplementary methods, Supplementary Material](#) online), and the GATA2-binding sites were identified by a pipeline based on genome-wide chromatin immunoprecipitation and sequencing (ChIP-seq) data from the ENCODE project (Dunham et al. 2012), as described separately (Arbiza L, Gronau I, Aksoy BA, Hubisz MJ, Gulko B, Keinan A, Siepel A, in revision). To improve efficiency, we performed the phylogenetic model fitting stage of our analysis in a preprocessing step. We fitted a neutral model estimated from 4-fold degenerate sites to the predesignated neutral sites by estimating two scale factors, one for the branch to the human genome (λ) and one for the other branches in the tree (λ^0 ; see Pollard et al. [2010] for details). This analysis assumed a ((human, chimpanzee), orangutan), rhesus macaque) tree topology. The outgroup scale, λ^0 , was estimated globally using all neutral sites genome-wide, and the human divergence scale, λ_b , was fitted separately in different genomic blocks. We used for this purpose a fixed set of 10 kb genomic windows overlapping by 5 kb and avoiding recombination hotspots. The same blocks were used for estimating the neutral polymorphism rates, θ_b . After fitting the phylogenetic model, we computed conditional distributions for the ancestral allele Z_i given the outgroup sequences at each nonfiltered nucleotide position i in the genome. The estimates of λ_b and θ_b , and the distributions for Z_i , were recorded in a database and used in all subsequent analyses.

Results

Simulations

We applied INSIGHT to various collections of synthetic elements to assess its accuracy and to validate our modeling assumptions. This was done by comparing our model-based parameter estimates both with “true” values reflecting the simulated evolutionary histories and with values obtained using simpler estimators based on counts of polymorphisms and divergences (see Materials and Methods). We simulated data sets roughly similar to our real data ([supplementary table S3, Supplementary Material](#) online), with 10,000–20,000 blocks each consisting of a 10 bp element flanked by 5 kbp of neutral sequence on each side. We considered a range of mixtures of neutral, weak negative (WN), strong negative (SN), and positive (P) selection (see Materials and Methods).

We started by examining four representative data sets (fig. 3A): 1) one with relatively few sites under selection (10%) and negative selection only (“Neg”); 2) another with a moderate fraction of sites under various types of selection (30%), including a substantial fraction under positive selection (“Pos”); 3) another with a high fraction of sites selection (60%), with mostly WN selection and no positive selection (“Weak”); and, finally, 4) a set with a substantial fraction of sites in each of the selective modes (“Mix”). We found that our model-based estimates were within one standard error of the true values across all mixtures of selective modes. The simple estimators also performed reasonably well in many cases, but the divergence-based estimators for ρ were strongly biased by positive selection (e.g., $\hat{\rho}_{\text{Div}} = -0.52$ in Pos and $\hat{\rho}_{\text{Div}} = -0.13$ in Mix). The reason for this bias is that these estimators implicitly attribute all divergence to neutral drift, an assumption that is violated by nonnegligible levels of positive selection. Similarly, the polymorphism-based estimator for ρ was biased downward in the presence of WN selection (e.g., $\hat{\rho}_{\text{Poly}} = 0.59$ and

$\rho_{\text{True}} = 0.8$ in Mix), because this estimator implicitly assumes that selection completely eliminates polymorphism, which is not true in this case. For similar reasons, the MK-based estimates of the number of adaptive divergences ($\hat{D}_{\text{p-MK}}$) were also biased in the presence of WN selection (Charlesworth and Eyre-Walker 2008).

These synthetic data sets—generated by forward simulation, under fairly realistic assumptions—also enabled us to directly evaluate the assumptions underlying our model. Consistent with our assumptions, no mutations reached fixation in the 340,000 negatively selected sites (weak or strong) in our synthetic data sets. On the other hand, polymorphisms under selection were not completely restricted to WN sites, as assumed; instead, 8% of them occurred in SN sites and 9% of them in positively selected sites, with the remaining 83% in WN sites. Nevertheless, our inference procedure appeared to be robust to these violations of our assumptions, with fairly accurate estimates of all parameters in all cases. Our default threshold of $f = 15\%$ for low-frequency polymorphisms appeared to be adequate: Only 4% of selected polymorphisms

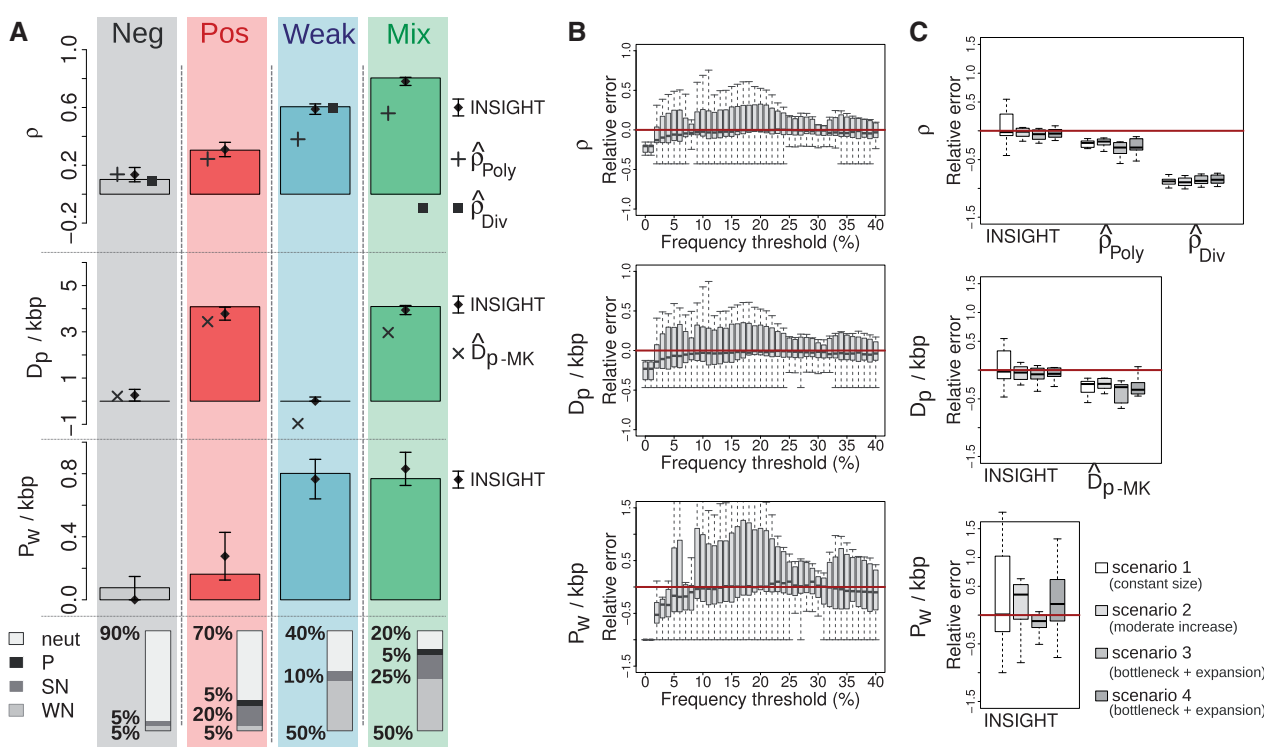


FIG. 3. Simulation results. (A) Parameter estimates for four collections of 20,000 simulated elements based on different mixtures of neutral (neut), positive (P), strong negative (SN), and weak negative (WN) selection (as indicated at bottom). The true values of ρ , D_p , and P_w are indicated by solid bars, and estimates from INSIGHT are indicated by diamonds, with error bars representing one standard error. For comparison, estimates from several simpler count-based methods are also shown, including estimates of ρ based on polymorphism ($\hat{\rho}_{\text{Poly}}$; “+”) and divergence ($\hat{\rho}_{\text{Div}}$; solid squares) rates, and estimates of D_p based on the MK framework ($\hat{D}_{\text{p-MK}}$; “x”). Adaptive divergences (D_p) and deleterious polymorphisms (P_w) are shown as rates per 1,000 base pairs (kbp). See Materials and Methods for details. (B) INSIGHT was applied to 11 collections of 10,000 elements with various fractions of sites under selection (see text), assuming a range of values for the low-frequency derived allele threshold f . Relative estimation errors for ρ , D_p , and P_w , measured as differences between the estimates and true values normalized by the true value, are shown as a function of the frequency threshold f . Each box plot describes the distribution of values for the 11 collections considered. Curvature-based standard errors for these experiments are summarize in [supplementary figure S2, Supplementary Material](#) online. (C) Simulated data sets were generated for the same 11 mixtures of selective modes \times four different demographic scenarios ([supplementary table S2, Supplementary Material](#) online), and INSIGHT parameter estimates were compared with true values. Box plots represent the distribution of relative error per demographic scenario. The relative estimation error for the simple count-based estimators, $\hat{\rho}_{\text{Poly}}$, $\hat{\rho}_{\text{Div}}$, and $\hat{D}_{\text{p-MK}}$, is shown for comparison.

exhibited derived allele frequencies exceeding this threshold, and these sites were vastly outnumbered by neutral high-frequency polymorphisms. Overall, although the simulated data did not fully support our modeling assumptions, only fairly minor violations were observed and our inference procedure seemed to be robust to them.

To test the robustness of INSIGHT to the assumed strength of selection, we conducted a series of simulations in which we perturbed the selection coefficients used for the three selective modes, SN, WN, and P (supplementary methods, Supplementary Material online). We found that INSIGHT performed very well across the entire range of values assumed for SN and positive selection (supplementary fig. S1, Supplementary Material online). A clear bias in parameter estimates was observed only when quite weak selection was assumed for the WN mode ($-2N_e s \leq 2$), in which case ρ and P_w were clearly underestimated and D_p was slightly underestimated. These results suggest that INSIGHT is generally robust to assumptions about selection coefficients, but that negative selection has to reach a certain threshold to produce sufficient shifts in derived allele frequencies to be detected by the method. As a result, our statements about ρ and P_w should be interpreted as including only the subset of deleterious polymorphisms that clear this threshold (supplementary methods, Supplementary Material online).

As mentioned earlier, $f = 15\%$ appears to be an adequate upper bound on the derived allele frequency at negatively selected sites (Fay et al. 2001; Zhang and Li 2005; Charlesworth and Eyre-Walker 2008). In reality, of course, this threshold depends on various factors, including the actual distribution of selection coefficients and the demographic history of the sample. To test the robustness of our model to the choice of f , we generated 11 collections of 10,000 elements with true fractions of sites under selection ranging from 0 to 1 (in steps of 0.1), keeping the proportion within selected sites in each collection constant at 45% WN, 50% SN, and 5% PD. We then applied INSIGHT to each data set using values of f ranging from 1% to 40% (fig. 1B). We found that very low thresholds ($f < 7\%$) resulted in clear underestimation of all model parameters, due to the presence of many selected polymorphisms with DAFs exceeding the cutoff, whereas very high thresholds ($f > 20\%$) led to high variance due to sparse data for high-frequency polymorphisms (supplementary fig. S2, Supplementary Material online). Importantly, however, no bias was observed for thresholds in the range of 7–20%, indicating robustness to the specific choice of threshold and justifying the default choice of 15%, which we used throughout most of our analysis.

An important feature of our model is that it directly contrasts sequence patterns in elements with those in nearby neutral sites, which should make it insensitive to the particular demographic history of the target population. To test robustness to demography, we simulated data sets for each of the 11 mixtures of selective modes described earlier using four different demographic scenarios for the target population: one with constant population size since divergence from chimpanzee, one with a moderate population expansion, and two others with a severe population bottleneck followed

by an exponential expansion (supplementary table S2, Supplementary Material online). Inference was performed separately for each of these 4×11 data sets, and the estimated parameters were then compared with their true values and with the simple count-based estimates (fig. 1C). The divergence-based estimates, $\hat{\rho}_{\text{Div}}$, were quite poor due to the effects of positive selection, as discussed earlier, and the polymorphism-based estimates, $\hat{\rho}_{\text{Poly}}$, consistently underestimated the true values, by an average of 20–30% across the different scenarios, due to the effects of WN selection. Similar patterns of underestimation were observed for the MK-based estimator of the number of adaptive divergences, $\hat{D}_{\text{p-MK}}$. In contrast, our model-based estimates of ρ and D_p showed no apparent bias in any of the simulated demographic scenarios. Estimates of the number of polymorphisms under selection, P_w , showed somewhat greater variance, as observed in our initial simulation study (fig. 1A), but the error in these estimates did not seem to be affected by demography. Thus, our method appears to be capable of disentangling the contributions of positive and negative selection even in the presence of a complex demographic history, without the need for explicit demographic inference.

Analysis of Human Noncoding Genomic Elements

To demonstrate its application to real data, we used INSIGHT to examine several classes of noncoding elements in the human genome, using 54 unrelated individual genomes from Complete Genomics to define human polymorphisms, and the chimpanzee, orangutan, and macaque genomes as outgroups (Materials and Methods). First, to assess our likelihood ratio cutoffs and ensure that our method adequately controls for false-positive inferences of selection, we applied INSIGHT to randomly selected “neutral” regions—excluding genes, conserved noncoding elements, and their immediate flanks (Materials and Methods). From the previously identified putatively neutral regions, we sampled 500 mutually exclusive collections of approximately 30,000 “neutral elements,” 10 bp long. For each collection, we estimated ρ and the corresponding LRT statistic for the null hypothesis of $\rho = 0$. The 500 estimated values of ρ were generally close to zero, with a median of 0.03 (supplementary fig. S3, Supplementary Material online) and almost no values > 0.1 . The distribution of LRT statistics was roughly similar to a 50:50 mixture of a point mass at zero and a χ^2 distribution with three degrees of freedom, as expected (Materials and Methods), but did show a clear shift toward large values relative to this distribution (fig. 4A). This shift may reflect violations of our simplifying assumptions in real genomic data (e.g., variation in mutation rates within blocks), contributions from alignment errors, or the inclusion of some functional sites within our “neutral” elements. Nevertheless, we found that the use of a more conservative (nonmixed) χ^2 distribution with three degrees of freedom adequately controlled for the excess in large LRT statistics. In particular, the empirical distribution of LRT statistics shows a good fit to the tail of this distribution (fig. 4A). Thus, we use this distribution for

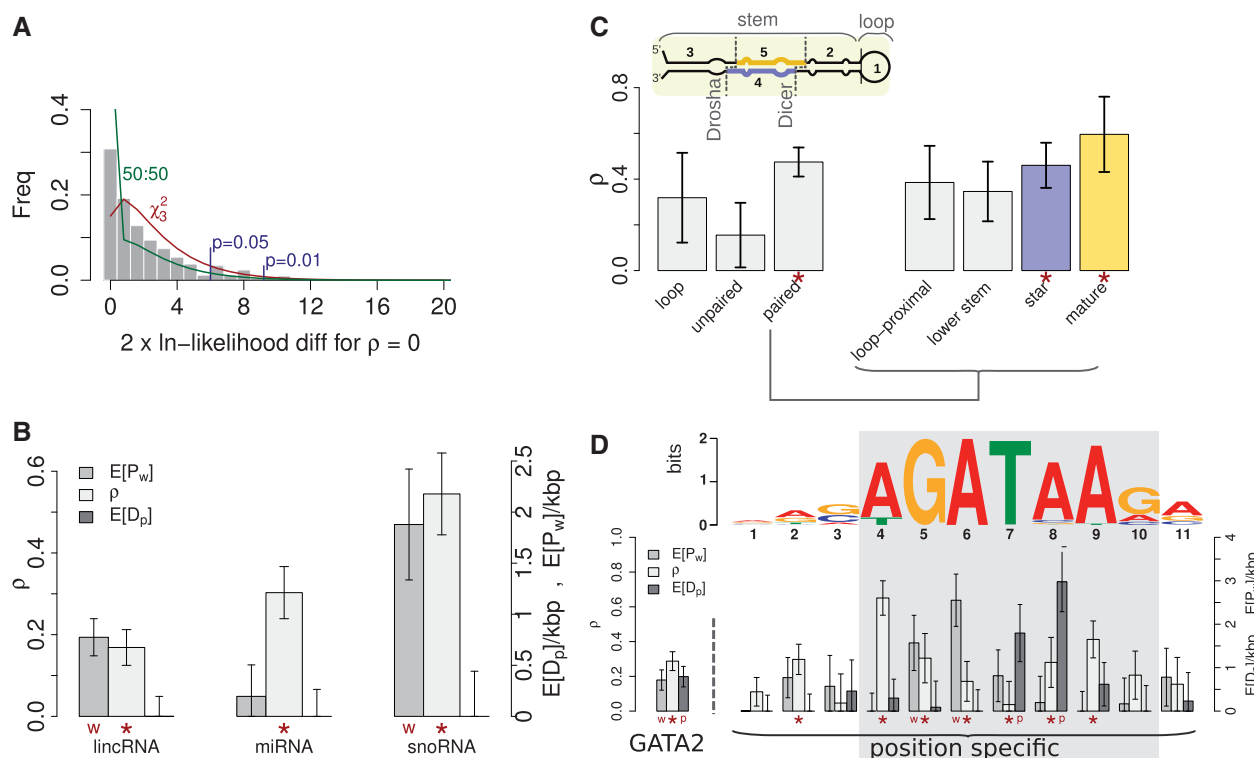


Fig. 4. Analysis of human genomic elements. (A) Distribution of LRT statistics for 500 sampled sets of “neutral” genomic elements, with approximately 30,000 elements per set. Test statistics reflect a null hypothesis that $\rho = 0$ and an alternative hypothesis that $\rho > 0$. For comparison, a χ^2_3 distribution (with three degrees of freedom; red) and a 50:50 mixture of a χ^2_3 distribution and a point mass at 0 (green) are also shown. Blue lines indicate significance thresholds for $P = 0.01$ and $P = 0.05$ based on the χ^2_3 distribution. Four of the 500 data sets (0.8%) had test statistics exceeding the $P = 0.01$ cutoff, and 24 (4.8%) exceeded the $P = 0.05$ cutoff, indicating a reasonably good fit to the tail of the distribution. The distribution of estimated values of ρ is shown in [supplementary figure S3, Supplementary Material](#) online. (B) Model-based estimates of ρ , $E[D_p]$, and $E[P_w]$ for three classes of noncoding RNAs (lincRNAs, miRNAs, and snoRNAs; see Materials and Methods). Error bars indicate one standard error. Symbols in red indicate statistical significance in LRTs for overall selection ($\rho > 0$; “*” $\rightarrow P < 0.01$) and WN selection ($\gamma > 0$; “w” $\rightarrow P < 0.01$), based on a χ^2_3 distribution for $\rho > 0$ and a χ^2_1 distribution for $\gamma > 0$. (C) Estimates of ρ for several structural regions of miRNAs (inset). (Left) Results for a coarse-grained partitioning into loop bases, unpaired stem bases, and paired stem bases. (Right) Results for a fine-grained partitioning of paired bases in the stem into loop-proximal, lower-stem, star, and mature regions, corresponding to the regions that undergo cropping and dicing by Drosha and Dicer (dashed lines). Estimates found to be significantly greater than 0 ($P \leq 0.01$) are highlighted (“*”). (D) The motif inferred for GATA2 together with position-specific estimates of ρ (left axis), D_p , and P_w (right axis). Statistical significance is assessed and indicated as in (B), with significant positive selection ($\eta > 0$; “P” $\rightarrow P < 0.01$) estimated using a χ^2_1 distribution. The “core” seven positions of the motif, having $IC > \frac{1}{2}$ are highlighted in gray. Estimates obtained for the joint analysis of all seven positions of the core motif are shown as well (left).

approximate calculations of nominal P values in our subsequent analyses.

Next, we examined three classes of noncoding RNAs annotated by the GENCODE project: microRNAs (miRNAs), small nucleolar RNAs (snoRNAs), and large interspersed noncoding RNAs (lincRNAs). We applied INSIGHT to a high-confidence subset of annotated elements in each of these five classes ([supplementary table S3 and methods, Supplementary Material](#) online). Our analysis considered various thresholds for distinguishing between low and high frequency polymorphisms, but our estimates were fairly insensitive to this threshold ([supplementary fig. S4, Supplementary Material](#) online), so we focus below on results for the default threshold of 15%. All three classes of elements were estimated to have significant fractions of sites under selection ($\rho > 0$; $P \leq 0.01$; [fig. 4B](#)). snoRNAs showed the highest estimated value ($\rho = 0.54 \pm 0.10$), consistent with their essential role in guiding chemical modifications of

ribosomal and transfer RNAs (Pang et al. 2006; Matera et al. 2007). miRNAs also showed a somewhat elevated estimate (0.3 ± 0.06). By contrast, lincRNAs were inferred to have a considerably smaller (but still significant) fraction of sites under selection (0.17 ± 0.04), consistent with previous observations indicating high levels of conservation are generally limited to short segments within lincRNAs (Guttman et al. 2009; Marques and Ponting 2009; Ulitsky et al. 2011). We also found significant evidence of WN selection in lincRNAs ($\gamma > 0$; $P \leq 0.01$). Furthermore, snoRNAs were estimated to have particularly high rates of weakly selected segregating polymorphisms ($E[P_w] = 1.9 \pm 0.5$ polymorphisms per kbp).

To shed additional light on the manner in which natural selection has influenced miRNA evolution, we applied INSIGHT separately to different structural components within the primary miRNA transcript ([fig. 4C, inset](#)). These structural classes were defined based on predictions of hairpin

secondary structures for the annotated miRNAs (supplementary methods, Supplementary Material online). We first partitioned the primary miRNA into loop and stem regions, distinguishing between paired and unpaired bases within the stem. Among these three partitions, paired bases in the stem region were estimated to have a particularly high fraction of sites under selection ($\rho = 0.47 \pm 0.06$; $P < 10^{-5}$; fig. 4C), consistent with their key role in stabilizing the hairpin structure. In contrast, the estimates for the other two classes were not found to be significantly greater than zero ($P > 0.05$). We further partitioned the stem into four subregions—loop-proximal, lower stem, star, and mature—reflecting the cleavage activity of Droscha and Dicer, the two RNase III cleavage enzymes of primary importance in miRNA biogenesis. We obtained the highest estimate of ρ (0.60 ± 0.16 ; $P \leq 0.01$; fig. 4C) for the 21–22 nt mature region, reflecting its dual role in structure preservation for efficient recognition and processing by Droscha and Dicer, and in sequence complementarity to target mRNAs. The lower-stem and loop-proximal regions had lower estimates of ρ , probably because they do not serve any direct regulatory role, but are important in preserving the hairpin structure. The star region had an intermediate estimate of ρ (0.46 ± 0.10 ; $P \leq 0.01$), perhaps because a fraction of star sequences are loaded into Argonaute complexes and carry out functional roles, even though most are degraded (Okamura et al. 2008).

Our estimates of recent selection in human miRNAs were generally concordant with previous comparative analyses in *Drosophila* based on patterns of divergence between species (Lai et al. 2003; Clark et al. 2007; Stark et al. 2007). They were also fairly consistent with estimates of sequence conservation across the primate phylogeny computed using phyloP scores (Pollard et al. 2010; supplementary fig. S5, Supplementary Material online). However, INSIGHT found somewhat stronger evidence for selection in the mature relative to the star region of the miRNA than did phyloP. This difference could reflect a shift toward WN selection in the star region, which is not apparent on comparative genomic time scales because selection is sufficiently strong to prohibit long-term fixation of derived alleles.

None of the analyzed noncoding RNAs showed strong evidence of positive selection, so we turned next to a collection of elements in which we expected to find evidence of adaptation based on other recent work (Arbiza L, Gronau I, Aksoy BA, Hubisz MJ, Gulko B, Keinan A, Siepel A, in revision): GATA2-binding sites. Using the available ChIP-seq-based annotations (Materials and Methods), we partitioned the nucleotides in binding sites into 11 groups, corresponding to the 11 positions in the GATA2 motif, and applied INSIGHT separately to each group (fig. 4D). We found that the signatures of natural selection were clearly concentrated in positions 4–10, which constitute the “core” region of the motif (fig. 4D). Indeed, six of these seven positions, and only one other position, were found to have significant estimates of ρ ($P \leq 0.01$). The 5th and 6th positions, associated with the “GA” portion of the GATA motif, were estimated to have undergone significant WN selection ($P < 0.01$), whereas the signature of positive selection came primarily from the 7th

and 8th positions, which are associated with the “TA” portion of the motif. Our posterior estimates indicated that nucleotides in these two positions contributed a total of 123 ± 43 adaptive divergences across approximately 27,500 binding sites. Interestingly, these positions (particularly the 8th) are known to play a role in modulating binding specificity of GATA2 (Ko and Engel 1993; Merika and Orkin 1993). They are also critical in determining the relative binding affinities of GATA1, GATA2, and GATA3, which regulate overlapping sets of genes and are known to serve as “switches” between alternative modes of gene expression (Bresnick et al. 2010; Dore et al. 2012).

Discussion

Recent advances in functional genomics have produced vast catalogs of candidate noncoding elements, including noncoding RNAs, transcription factor binding sites, RNA-binding sites, and many others (Gerstein et al. 2010; Dunham et al. 2012). Evolutionary analyses will be essential in improving annotations of these elements and revealing their functional roles. Although methods based on patterns of divergence between species have become widely used in the study of noncoding functional elements, these methods are limited by their consideration of relatively long evolutionary time scales and their sensitivity to alignment errors and other technical artifacts. INSIGHT is intended to complement these methods by detecting signatures of recent selection based on newly available population genomic data and comparative genomic data for closely related species.

INSIGHT bears some similarities to MK-based methods (McDonald and Kreitman 1991; Smith and Eyre-Walker 2002; Bierne and Eyre-Walker 2004; Andolfatto 2005), Poisson random field (PRF)-based methods (Sawyer and Hartl 1992; Bustamante et al. 2002, 2005; Williamson et al. 2005), and related methods for characterizing the DFE (Eyre-Walker et al. 2006; Boyko et al. 2008; Eyre-Walker and Keightley 2009), but it differs from previous methods in several important respects. Unlike MK-based methods, INSIGHT is based on a full generative probabilistic model, it explicitly models WN selection, and it pools information from many loci in a manner that properly accommodates differences in mutation rates and coalescence time across the genome. Unlike PRF-based methods, it does not attempt to directly model the complex relationship between allele frequencies and natural selection, but instead works by contrasting patterns of polymorphism and divergence in the elements of interest and flanking sites. INSIGHT additionally allows for straightforward LRTs of various hypotheses of interest, and it allows parameter variances to be approximately characterized using standard methods. For these reasons, we expect it to be a valuable addition to the arsenal of methods available for analysis of polymorphism and divergence data.

The INSIGHT model is designed to exploit newly available genome-scale data sets describing both candidate functional elements (Gerstein et al. 2010; Roy et al. 2010; Dunham et al. 2012) and variation within populations (1000 Genomes Project Consortium 2010; Mackay et al. 2012). Although the underlying graphical model is not complex (fig. 2), a naive

approach to parameter estimation would still be prohibitively CPU-intensive with genome-wide data. We achieve major gains in efficiency by decomposing the inference procedure into three separate steps, concerned with the estimation of the phylogenetic, neutral, and selection parameters, respectively. This decomposition relies on the simplifying assumption that sites within the elements of interest contain negligible information about the neutral parameters of the model, because they are vastly outnumbered by the flanking neutral sites—a property that can typically be guaranteed by construction. It also depends on the use of a single phylogenetic model per locus in estimating the prior distribution of the ancestral allele at all sites, which should be adequate as long as relatively close outgroups are used. Notably, the first two of these steps can be performed in preprocessing and reused in the analysis of any set of loci that use the same flanking regions. Furthermore, the neutral flanks can be designed to maximize the potential for reuse (as we have done here) by defining a set of fixed genomic blocks, and associating each element with the neutral sites of the nearest block. This strategy allows the neutral and phylogenetic parameters to be pre-estimated for each block and reused in any number of subsequent analyses. Importantly, these steps dominate the running time of the inference algorithm (particularly the phylogenetic estimation step). The final stage, in which the parameters ρ , η , and γ are estimated, is independent of the number of genomes considered and typically takes less than a minute.

It is worth emphasizing that INSIGHT can be applied to any collection of genomic elements, provided each one is sufficiently short that it does not span regions having markedly different mutation rates or coalescence times, and provided each element can be associated with nearby sites likely to be free from the effects of selection. In this article, we have focused on the case of genome-wide collections of elements of a particular type, such as miRNAs or binding sites for a particular transcription factor, but many other types of analysis are possible. For example, in related work (Arbiza L, Gronau I, Aksoy BA, Hubisz MJ, Gulko B, Keinan A, Siepel A, in revision), we have examined various subsets of transcription factor binding sites, such as those associated with genes of a particular Gene Ontology category or expressed at a particular level, and those having various levels of predicted binding affinity. As we have shown here, the method can also be used to pinpoint signals of selection within elements, by partitioning them based on structural features (e.g., particular motif positions or particular miRNA structural compartments). Similar analyses could be used to contrast regions of the genome having different epigenomic marks, sequences nearby and far from genes, **sequences on sex chromosomes and autosomes**, or any number of other biologically significant genomic partitions.

INSIGHT could be extended in various ways to improve the fit of the model to the data and broaden the utility of the program. In this analysis, we had a sufficiently complete collection of human variation data to simply discard positions with missing data in one or more samples. In cases of more missing data, however, it may be worthwhile to use the

strategy of adjusting Watterson's constant a_n in the appropriate conditional distributions (table 3) based on the number of samples for which data are available at each genomic position. This simple approach should work well as long as the amount of missing data is not excessive, but it will require some care in programming to efficiently accommodate site-wise variation in a_n (supplementary methods, Supplementary Material online). Another useful extension would be to allow for variation across loci in the global parameters ρ , η , and γ , say, by assuming locus-specific parameters are drawn from (discretized) Beta (for ρ) or Gamma (for η and γ) distributions and estimating the hyper-parameters for these distributions from the data. This strategy should improve model fit considerably in cases of variable selection across loci, similar to phylogenetic models that allow for rate variation among sites (Yang 1994). A further extension would be to use a fully Bayesian approach and infer posterior distributions for the parameters of interest. This would also be fairly straightforward, but would most likely require Markov chain Monte Carlo sampling or variational Bayes approximations. These and other extensions would help further in using patterns of polymorphism and divergence to shed light on recent evolutionary processes, particularly in noncoding regions, and may improve predictions of the fitness effects of mutations across the genome.

Supplementary Material

Supplementary methods, tables S1–S3, and figures S1–S5 are available at *Molecular Biology and Evolution* online (<http://www.mbe.oxfordjournals.org/>).

Acknowledgments

This research was supported by a David and Lucile Packard Fellowship for Science and Engineering, National Science Foundation grant DBI-0644111, and National Institutes of Health (NIGMS) grant GM102192. Additional support was provided by a postdoctoral fellowship from the Cornell Center for Vertebrate Genomics (to L.A.).

References

- Andolfatto P. 2005. Adaptive evolution of non-coding DNA in *Drosophila*. *Nature* 437:1149–1152.
- Bierne N, Eyre-Walker A. 2004. The genomic rate of adaptive amino acid substitution in *Drosophila*. *Mol Biol Evol.* 21:1350–1360.
- Boffelli D, McAuliffe J, Ovcharenko D, Lewis KD, Ovcharenko I, Pachter L, Rubin EM. 2003. Phylogenetic shadowing of primate sequences to find functional regions of the human genome. *Science* 299:1391–1394.
- Boyko AR, Williamson SH, Indap AR, et al. (14 co-authors). 2008. Assessing the evolutionary impact of amino acid mutations in the human genome. *PLoS Genet.* 4:e1000083.
- Bresnick EH, Lee HY, Fujiwara T, Johnson KD, Keles S. 2010. GATA switches as developmental drivers. *J Biol Chem.* 285:31087–31093.
- Bustamante CD, Fledel-Alon A, Williamson S, et al. (11 co-authors). 2005. Natural selection on protein-coding genes in the human genome. *Nature* 437:1153–1157.
- Bustamante CD, Nielsen R, Sawyer SA, Olsen KM, Purugganan MD, Hartl DL. 2002. The cost of inbreeding in *Arabidopsis*. *Nature* 416:531–534.
- Charlesworth J, Eyre-Walker A. 2008. The McDonald-Kreitman test and slightly deleterious mutations. *Mol Biol Evol.* 25:1007–1015.

- Chernoff H. 1954. On the distribution of the likelihood ratio. *Ann Math Stat.* 25:573–578.
- Clark A, Eisen M, Smith D, et al. (11 co-authors). 2007. Evolution of genes and genomes on the *Drosophila* phylogeny. *Nature* 450:203–218.
- Cooper GM, Stone EA, Asimenos G, Green ED, Batzoglou S, Sidow A. 2005. Distribution and intensity of constraint in mammalian genomic sequence. *Genome Res.* 15:901–913.
- Dermitzakis ET, Clark AG. 2002. Evolution of transcription factor binding sites in mammalian gene regulatory regions: conservation and turnover. *Mol Biol Evol.* 19:1114–1121.
- Dore LC, Chlon TM, Brown CD, White KP, Crispino JD. 2012. Chromatin occupancy analysis reveals genome-wide GATA factor switching during hematopoiesis. *Blood* 119:3724–3733.
- Drmanac R, Sparks AB, Callow MJ, et al. (65 co-authors). 2010. Human genome sequencing using unchained base reads on self-assembling DNA nanoarrays. *Science* 327:78–81.
- Dunham I, Kundaje A, Aldred SF, et al. (1196 co-authors). 2012. An integrated encyclopedia of DNA elements in the human genome. *Nature* 489:57–74.
- Eyre-Walker A, Keightley PD. 2009. Estimating the rate of adaptive molecular evolution in the presence of slightly deleterious mutations and population size change. *Mol Biol Evol.* 26:2097–2108.
- Eyre-Walker A, Woolfit M, Phelps T. 2006. The distribution of fitness effects of new deleterious amino acid mutations in humans. *Genetics* 173:891–900.
- Fay JC, Wyckoff GJ, Wu CI. 2001. Positive and negative selection on the human genome. *Genetics* 158:1227–1234.
- Gerstein MB, Lu ZJ, Van Nostrand EL, et al. (131 co-authors). 2010. Integrative analysis of the *Caenorhabditis elegans* genome by the modENCODE project. *Science* 330:1775–1787.
- Guttman M, Amit I, Garber M, et al. (20 co-authors). 2009. Chromatin signature reveals over a thousand highly conserved large non-coding RNAs in mammals. *Nature* 458:223–227.
- Harrow J, Frankish A, Gonzalez JM, et al. (41 co-authors). 2012. GENCODE: the reference human genome annotation for The ENCODE Project. *Genome Res.* 22:1760–1774.
- Hernandez RD. 2008. A flexible forward simulator for populations subject to selection and demography. *Bioinformatics* 24:2786–2787.
- Hubisz MJ, Pollard KS, Siepel A. 2011. PHAST and RPHAST: phylogenetic analysis with space/time models. *Brief Bioinform.* 12:41–51.
- Ko LJ, Engel JD. 1993. DNA-binding specificities of the GATA transcription factor family. *Mol Cell Biol.* 13:4011–4022.
- Kondrashov AS, Crow JF. 1993. A molecular approach to estimating the human deleterious mutation rate. *Hum Mutat.* 2:229–234.
- Lai EC, Tomancak P, Williams RW, Rubin GM. 2003. Computational identification of *Drosophila* microRNA genes. *Genome Biol.* 4:R42.
- Lehmann EEL, Casella G. 1998. Theory of point estimation. New York: Springer.
- Mackay TF, Richards S, Stone EA, et al. (52 co-authors). 2012. The *Drosophila melanogaster* genetic reference panel. *Nature* 482:173–178.
- Marques AC, Ponting CP. 2009. Catalogues of mammalian long non-coding RNAs: modest conservation and incompleteness. *Genome Biol.* 10:R124.
- Matera AG, Terns RM, Terns MP. 2007. Non-coding RNAs: lessons from the small nuclear and small nucleolar RNAs. *Nat Rev Mol Cell Biol.* 8:209–220.
- McDonald JH, Kreitman M. 1991. Adaptive protein evolution at the Adh locus in *Drosophila*. *Nature* 351:652–654.
- Merika M, Orkin SH. 1993. DNA-binding specificity of GATA family transcription factors. *Mol Cell Biol.* 13:3999–4010.
- Moses AM, Pollard DA, Nix DA, Iyer VN, Li XY, Biggin MD, Eisen MB. 2006. Large-scale turnover of functional transcription factor binding sites in *Drosophila*. *PLoS Comput Biol.* 2:e130.
- Nielsen R, Hellmann I, Hubisz M, Bustamante C, Clark A. 2007. Recent and ongoing selection in the human genome. *Nat Rev Genet.* 8:857–868.
- Okamura K, Phillips MD, Tyler DM, Duan H, Chou YT, Lai EC. 2008. The regulatory activity of microRNA* species has substantial influence on microRNA and 3' UTR evolution. *Nat Struct Mol Biol.* 15:354–363.
- 1000 Genomes Project Consortium. 2010. A map of human genome variation from population-scale sequencing. *Nature* 467:1061–1073.
- Pang KC, Frith MC, Mattick JS. 2006. Rapid evolution of noncoding RNAs: lack of conservation does not mean lack of function. *Trends Genet.* 22:1–5.
- Pollard KS, Hubisz MJ, Rosenbloom KR, Siepel A. 2010. Detection of nonneutral substitution rates on mammalian phylogenies. *Genome Res.* 20:110–121.
- Roy S, Ernst J, Kharchenko PV, et al. (269 co-authors). 2010. Identification of functional elements and regulatory circuits by *Drosophila* modENCODE. *Science* 330:1787–1797.
- Sawyer SA, Hartl DL. 1992. Population genetics of polymorphism and divergence. *Genetics* 132:1161–1176.
- Self S, Liang K. 1987. Asymptotic properties of maximum likelihood estimators and likelihood ratio tests under nonstandard conditions. *J Am Stat Assoc.* 82:605–610.
- Siepel A, Bejerano G, Pedersen JS, et al. (16 co-authors). 2005. Evolutionarily conserved elements in vertebrate, insect, worm, and yeast genomes. *Genome Res.* 15:1034–1050.
- Smith NG, Eyre-Walker A. 2002. Adaptive protein evolution in *Drosophila*. *Nature* 415:1022–1024.
- Stark A, Kheradpour P, Parts L, Brennecke J, Hodges E, Hannon GJ, Kellis M. 2007. Systematic discovery and characterization of fly microRNAs using 12 *Drosophila* genomes. *Genome Res.* 17:1865–1879.
- Stoletzki N, Eyre-Walker A. 2011. Estimation of the neutrality index. *Mol Biol Evol.* 28:63–70.
- Thomas JW, Touchman JW, Blakesley RW, et al. (70 co-authors). 2003. Comparative analyses of multi-species sequences from targeted genomic regions. *Nature* 424:788–793.
- Ulitsky I, Shkumatava A, Jan CH, Sive H, Bartel DP. 2011. Conserved function of lincRNAs in vertebrate embryonic development despite rapid sequence evolution. *Cell* 147:1537–1550.
- Watterson GA. 1975. On the number of segregating sites in genetical models without recombination. *Theor Popul Biol.* 7:256–276.
- Williamson SH, Hernandez R, Fledel-Alon A, Zhu L, Nielsen R, Bustamante CD. 2005. Simultaneous inference of selection and population growth from patterns of variation in the human genome. *Proc Natl Acad Sci U S A.* 102:7882–7887.
- Wilson DJ, Hernandez RD, Andolfatto P, Przeworski M. 2011. A population genetics-phylogenetics approach to inferring natural selection in coding sequences. *PLoS Genet.* 7:e1002395.
- Yang Z. 1994. Maximum likelihood phylogenetic estimation from DNA sequences with variable rates over sites: approximate methods. *J Mol Evol.* 39:306–314.
- Yi X, Liang Y, Huerta-Sanchez E, et al. (70 co-authors). 2010. Sequencing of 50 human exomes reveals adaptation to high altitude. *Science* 329:75–78.
- Zhang L, Li WH. 2005. Human SNPs reveal no evidence of frequent positive selection. *Mol Biol Evol.* 22:2504–2507.